

构建统一的大数据分析+AI 流水

史栋杰· Intel / 资深软件架构
师

Apache Flink Community China

CONTENT

目录 >>

01 /
Overview

02 /
BigDL

03 /
Analytics Zoo

04 /
Analytics Zoo
Inference on Flink

01

Overview

AI on Big Data



Distributed, High-Performance
Deep Learning Framework

<https://github.com/intel-analytics/bigdl>

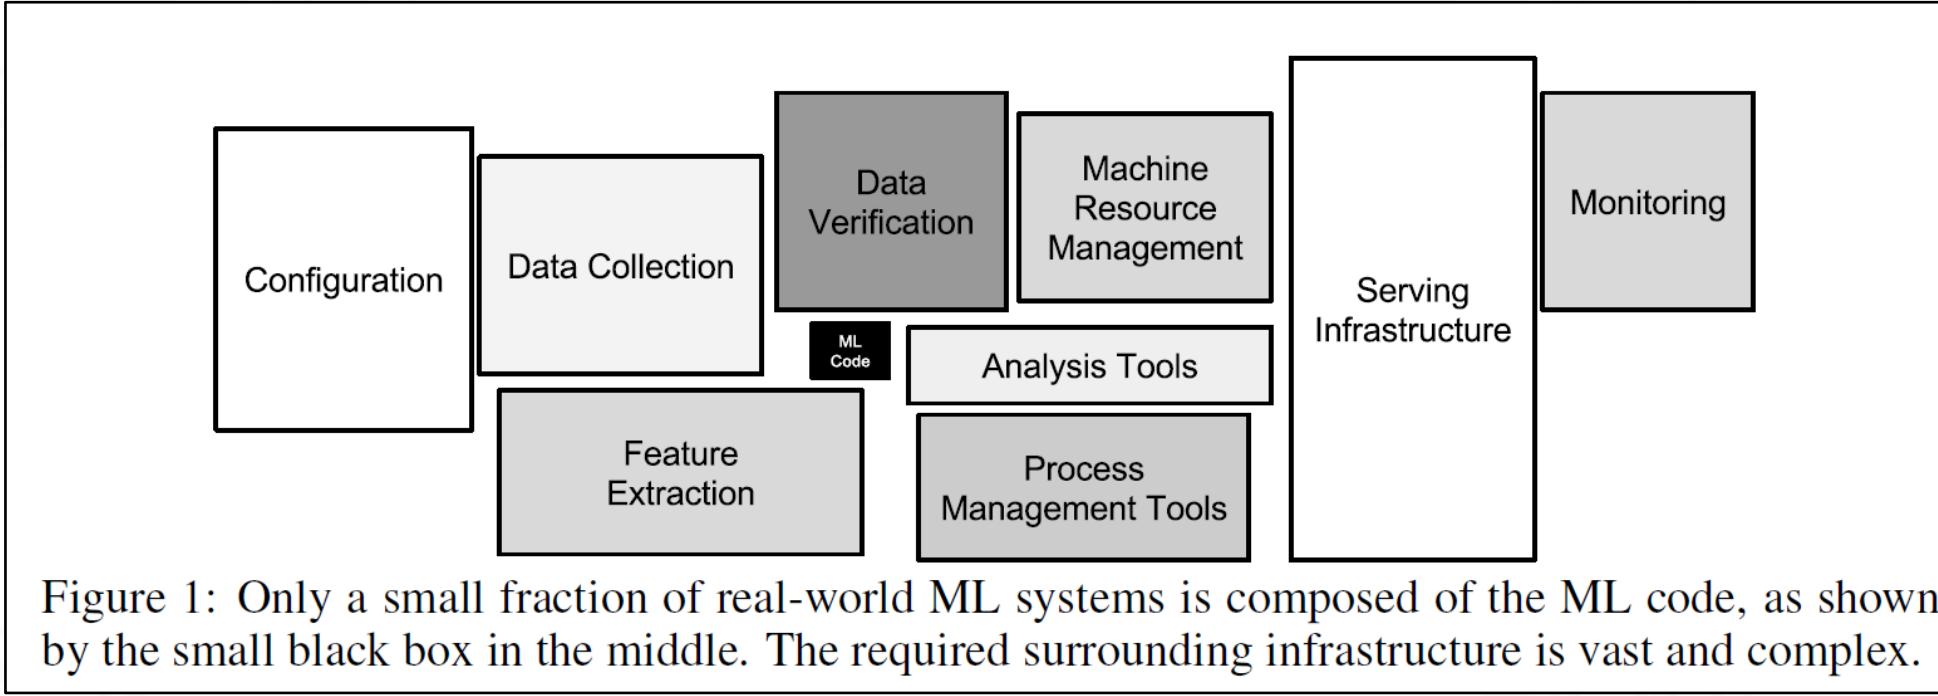


Analytics + AI Platform
Distributed TensorFlow*, Keras*, PyTorch* and BigDL

<https://github.com/intel-analytics/analytics-zoo>

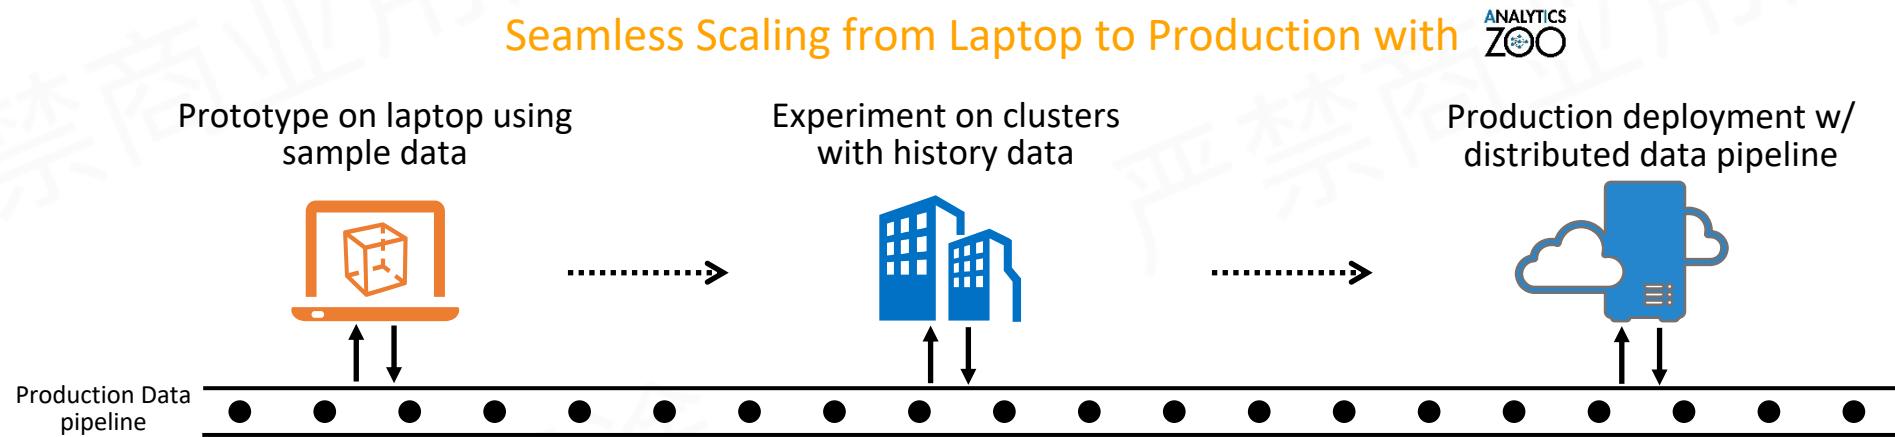
Accelerating Data Analytics + AI Solutions At Scale

Real-World ML/DL Applications Are Complex Data Analytics Pipelines



“Hidden Technical Debt in Machine Learning Systems”,
Sculley et al., Google, NIPS 2015 Paper

End-to-End Big Data Analytics and AI Pipeline



- “Zero” code change from laptop to distributed cluster
- Directly access production data (Hadoop/Hive/HBase) without data copy
- Easily prototype the end-to-end pipeline
- Seamlessly deployed on production big data clusters

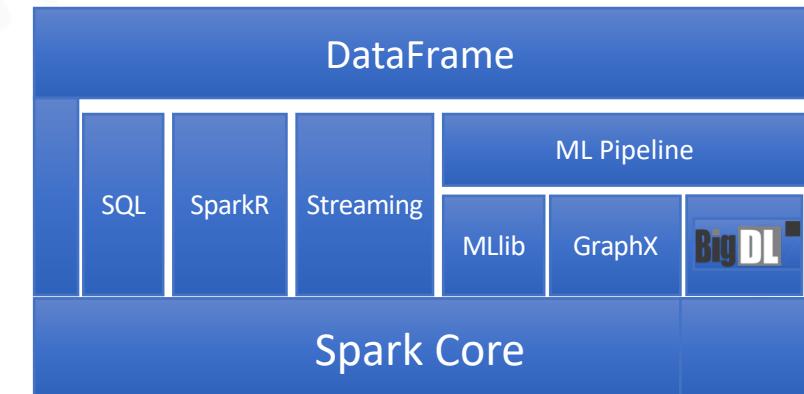
02

BigDL

BigDL

Bringing Deep Learning To Big Data Platform

- **Distributed** deep learning framework for Apache Spark
- Make deep learning more accessible to **big data users** and **data scientists**
 - Write deep learning applications as **standard Spark programs**
 - Run on existing Spark/Hadoop clusters (**no changes needed**)
- Feature parity with popular deep learning frameworks
 - E.g., Caffe, Torch, Tensorflow, etc.
- High performance (on CPU)
 - Powered by Intel MKL and multi-threaded programming
- Efficient scale-out
 - Leveraging Spark for distributed training & inference



<https://github.com/intel-analytics/BigDL>

<https://bigdl-project.github.io/>

BigDL Run as Standard Spark Programs

Standard Spark jobs

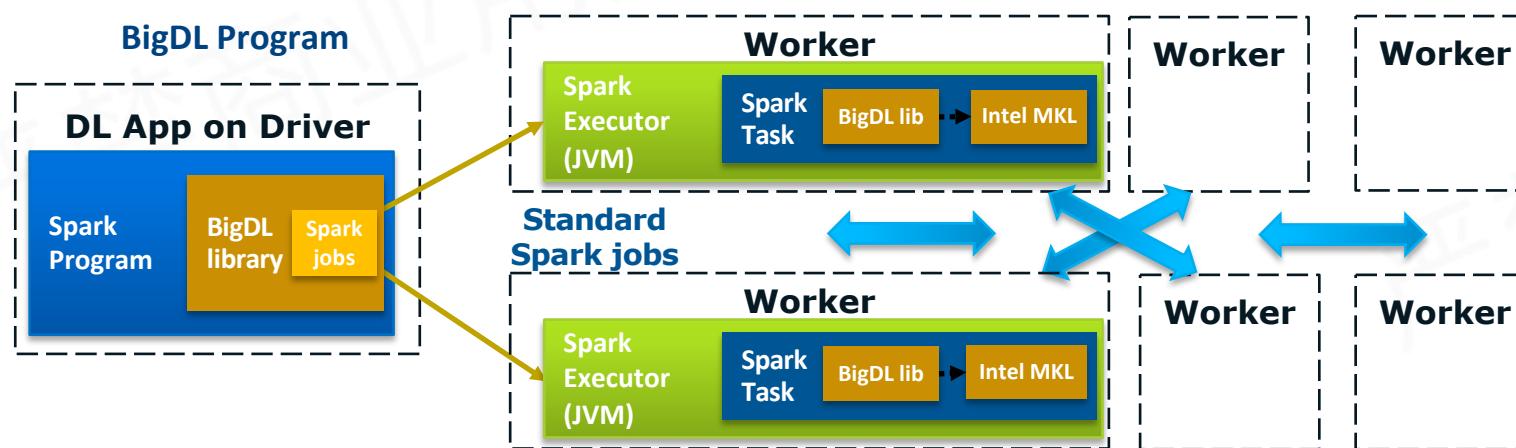
- No changes to the Spark or Hadoop clusters needed

Iterative

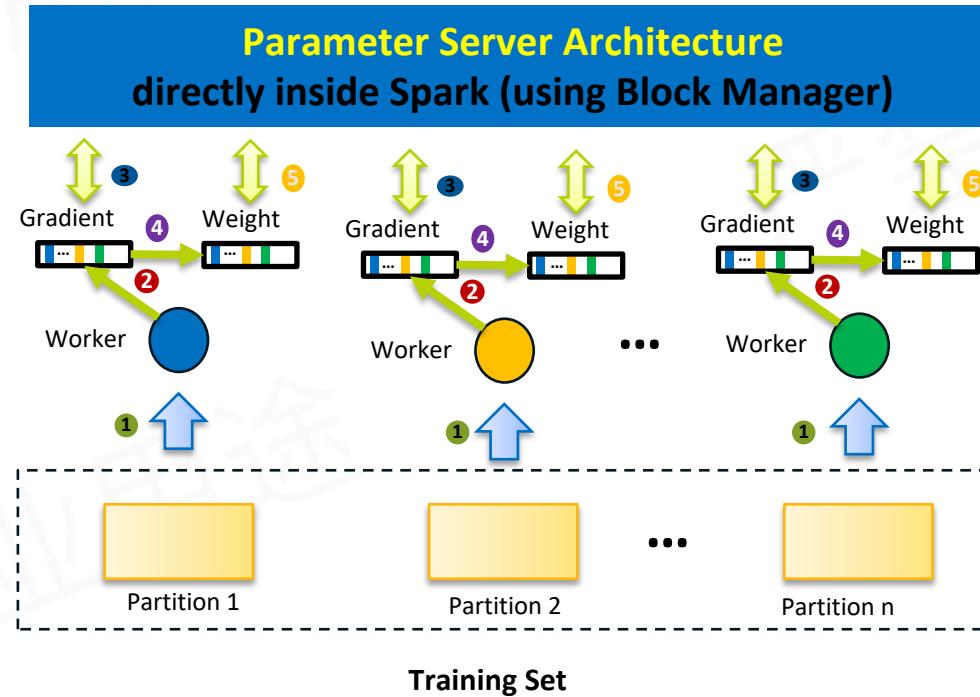
- Each iteration of the training runs as a Spark job

Data parallel

- Each Spark task runs the same model on a subset of the data (batch)



Distributed Training in BigDL



Peer-2-Peer All-Reduce synchronization

03

Analytics Zoo

Analytics Zoo

Unified Analytics + AI Platform for Big Data

Use case	Recommendation	Anomaly Detection	Text Classification	Text Matching		
Model	Image Classification	Object Detection	Seq2Seq	Transformer	BERT	
Feature Engineering	image	3D image	text	Time series		
High Level Pipelines	tfpark: Distributed TF on Spark	Distributed Keras w/ autograd on Spark	nnframes: Spark Dataframes & ML Pipelines for Deep Learning	Distributed Model Serving (batch, streaming & online)		
Backend/ Library	TensorFlow	Keras	BigDL	NLP Architect	Apache Spark	Apache Flink
	MKLDNN	OpenVINO	Intel® Optane™ DCPMM	DL Boost (VNNI)		

<https://github.com/intel-analytics/analytics-zoo>

Analytics Zoo

Unified Analytics + AI Platform for Big Data

Build end-to-end deep learning applications for big data

- Distributed *TensorFlow* on Spark
- *Keras* API (with autograd & transfer learning support) on Spark
- *nnframes*: native DL support for Spark DataFrames and ML Pipelines

Productionize deep learning applications for big data at scale

- Plain Java/Python *model serving* APIs (w/ OpenVINO support)
- Support Web Services, Spark, Flink, Storm, Kafka, etc.

Out-of-the-box solutions

- Built-in deep learning *models*, *feature engineering* operations, and reference *use cases*

Distributed TF & Keras on Spark

- Data wrangling and analysis using PySpark
- Deep learning model development using TensorFlow or Keras
- Distributed training / inference on Spark

Write TensorFlow code inline in PySpark program

```
#pyspark code
train_rdd = spark.hadoopFile(...).map(...)
dataset = TFDataSet.from_rdd(train_rdd,...)

#tensorflow code
import tensorflow as tf
slim = tf.contrib.slim
images, labels = dataset.tensors
with slim.arg_scope(lenet.lenet_arg_scope()):
    logits, end_points = lenet.lenet(images, ...)
loss = tf.reduce_mean( \
    tf.losses.sparse_softmax_cross_entropy( \
        logits=logits, labels=labels))

#distributed training on Spark
optimizer = TFOptimizer.from_loss(loss, Adam(...))
optimizer.optimize(end_trigger=MaxEpoch(5))
```

Spark Dataframe & ML Pipeline for DL

```
#Spark dataframe transformations
parquetfile = spark.read.parquet(...)
train_df = parquetfile.withColumn(...)

#Keras API
model = Sequential()
    .add(Convolution2D(32, 3, 3, activation='relu', input_shape=...)) \
    .add(MaxPooling2D(pool_size=(2, 2))) \
    .add(Flatten()).add(Dense(10, activation='softmax')))

#Spark ML pipeline
Estimator = NNEstimator(model, CrossEntropyCriterion()) \
    .setLearningRate(0.003).setBatchSize(40).setMaxEpoch(5) \
    .setFeaturesCol("image")
nnModel = estimator.fit(train_df)
```

Distributed Model Serving

Overview Timeline Exceptions Configuration

+ -

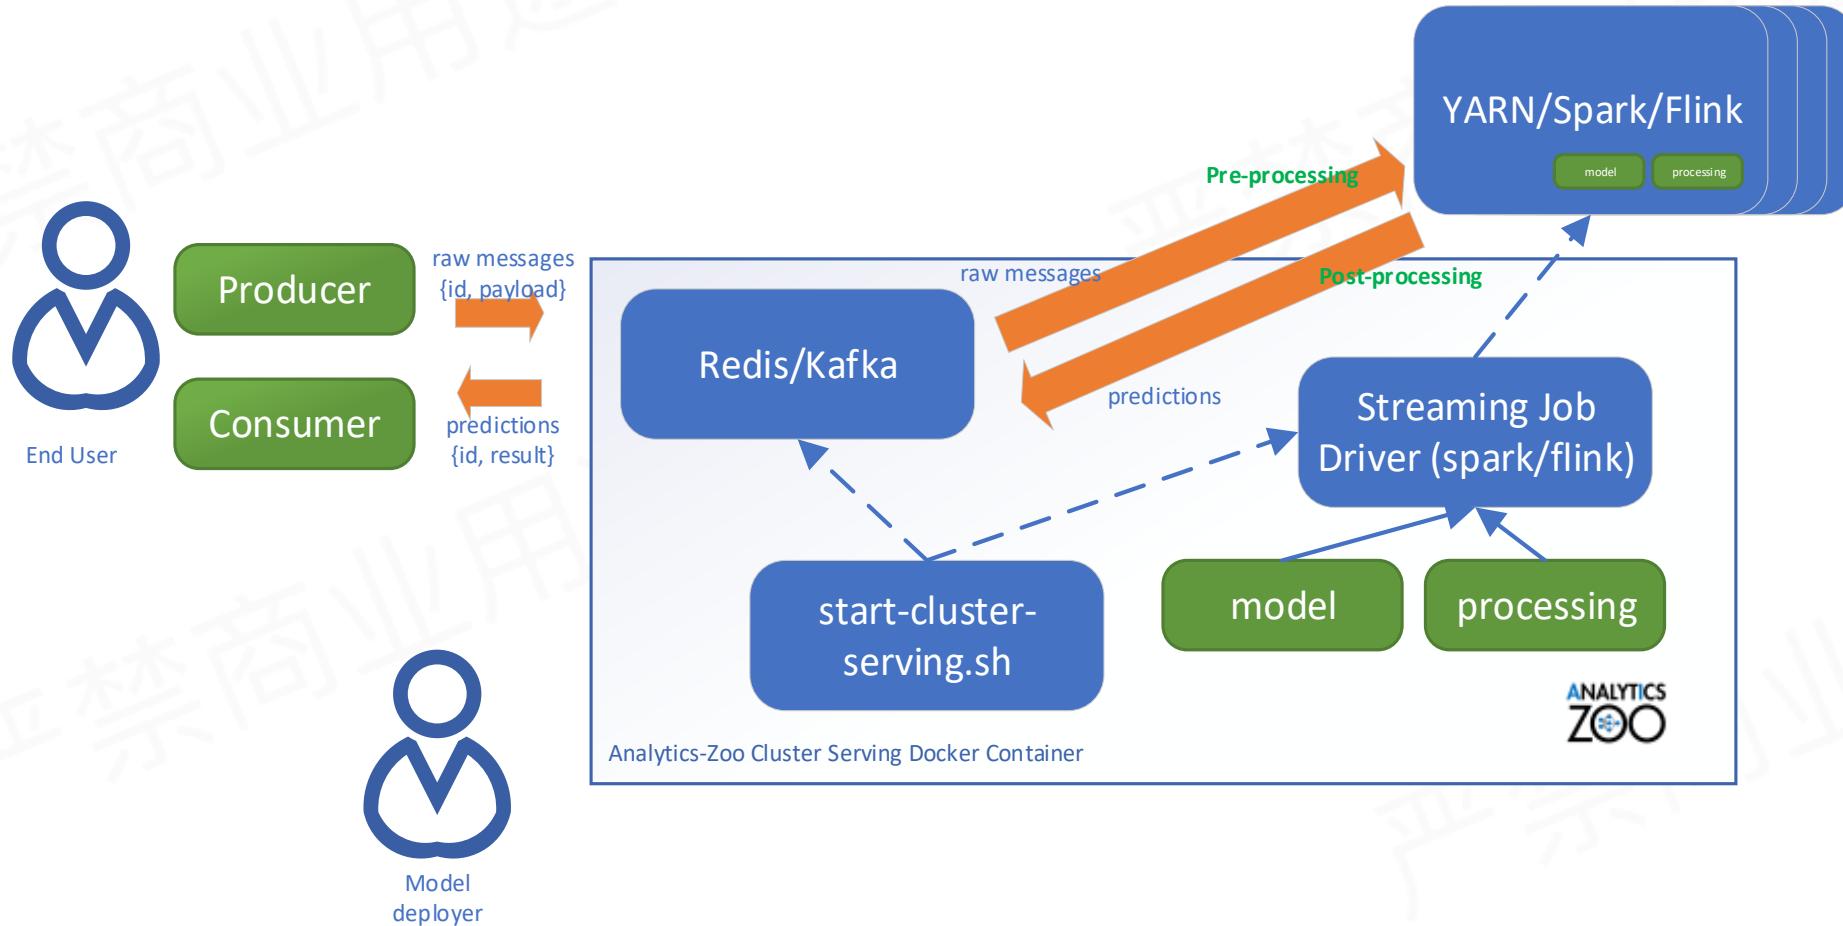
Analytics Zoo Model

Aggregate task status by Task ID / Day										
Start Time	End Time	Duration	Name	Bytes received	Records received	Bytes sent	Records sent	Parallelism	Tasks	Status
2019-08-01, 10:09:36	2019-08-01, 10:13:20	3m 44s	Source: Custom Source	0 B	0	537 MB	330	1	0 0 1 0 0 0 0	RUNNING
2019-08-01, 10:09:36	2019-08-01, 10:13:20	3m 44s	Flat Map -> Map	536 MB	326	0 B	320	6	0 0 6 0 0 0 0	RUNNING
Start Time	End Time	Duration	Bytes received	Records received	Bytes sent	Records sent	Attempt	Host	Status	
2019-08-01, 10:09:40		3m 40s	83.4 MB	53	0 B	52	1	Almaren-Node-018:41072	RUNNING	
2019-08-01, 10:09:40		3m 40s	80.4 MB	54	0 B	53	1	Almaren-Node-018:41072	RUNNING	
2019-08-01, 10:09:36		3m 44s	76.2 MB	54	0 B	53	1	Almaren-Node-019:39744	RUNNING	
2019-08-01, 10:09:40		3m 40s	96.9 MB	55	0 B	54	1	Almaren-Node-019:39744	RUNNING	
2019-08-01, 10:09:40		3m 40s	87.8 MB	55	0 B	54	1	Almaren-Node-020:35343	RUNNING	

Distributed model serving in Web Service, Flink, Kafka, Storm, etc.

- Plain Java or Python API, with OpenVINO and DL Boost (VNNI) support

Distributed Model Serving



OpenVINO Support for Model Serving

```
from zoo.common.nncontext import init_nncontext
from zoo.feature.image import ImageSet
from zoo.pipeline.inference import InferenceModel

sc = init_nncontext("OpenVINO Object Detection Inference Example")
images = ImageSet.read(options.img_path, sc,
                      resize_height=600, resize_width=600).get_image().collect()
input_data = np.concatenate([image.reshape((1, 1) + image.shape) for image in images], axis=0)

model = InferenceModel()
model.load_tf(options.model_path, backend="openvino", model_type=options.model_type)
predictions = model.predict(input_data)

# Print the detection result of the first image.
print(predictions[0])
```

Transparently support OpenVINO in model serving, which deliver a significant boost for inference speed

Upcoming Analytics Zoo 0.6 Release

- **Distributed PyTorch on Spark**
- **Ray on Spark**
 - Run Ray programs directly on standard Hadoop/YARN clusters
- **AutoML support**
 - Automatic feature generation, model selection and hyper-parameter tuning for *time series prediction*
- **Cluster serving**
 - Distributed, real-time (streaming) model serving with simple pub-sub interface

Use Cases

Technology



bluedata®

cloudera®



THE SUPERCOMPUTER COMPANY



DELL EMC

inspur



innovate with confidence



Qubole

Cloud Service Providers



Alibaba Cloud
aliyun.com



Azure

Tencent 腾讯



IBM Cloud



End Users

cdhi

Telefónica



中国电信
CHINA TELECOM



THE WORLD BANK

JD.com

CERN openlab

Midea®

韵达 · 大 ·
EXPRESS

小米
CISCO™

UnionPay
银联

software.intel.com/AlonBigData

Not a full list

*Other names and brands may be claimed as the property of others.

04

Analytics Zoo Inference on Flink

Analytics Zoo Inference Model

- **Java AbstractInferenceModel**

- public void load(String modelPath, String weightPath)
- public void loadCaffe(String modelPath, String weightPath)
- public void loadTF(String modelPath)
- public void loadTF(String modelPath, ...)
- public void loadTF(String savedModelDir, ...)
- public void loadTF(byte[] savedModelBytes, ...)
- public void loadOpenVINO(String modelFilePath, String weightFilePath...)
- public List<List<JTensor>> predict(List<List<JTensor>> inputs)

- **Scala InferenceModel**

- **POJO style**
- **Thread-Safe**

pom

```
<repositories>
  <repository>
    <id>ossrh</id>
    <name>sonatype repository</name>
    <url>https://oss.sonatype.org/content/groups/public/</url>
    <releases>
      <enabled>true</enabled>
    </releases>
    <snapshots>
      <enabled>true</enabled>
    </snapshots>
  </repository>
</repositories>
```

```
<dependency>
  <groupId>com.intel.analytics.zoo</groupId>
  <artifactId>analytics-zoo-bigdl_0.9.0-spark_2.4.3</artifactId>
  <version>0.6.0-SNAPSHOT</version>
</dependency>
<dependency>
  <groupId>com.intel.analytics.zoo</groupId>
  <artifactId>zoo-core-dist-all</artifactId>
  <version>0.6.0-SNAPSHOT</version>
  <type>jar</type>
</dependency>
<dependency>
  <groupId>org.scala-lang</groupId>
  <artifactId>scala-library</artifactId>
  <version>2.11.11</version>
</dependency>
<dependency>
  <groupId>org.apache.flink</groupId>
  <artifactId>flink-scala_2.11</artifactId>
  <version>1.8.1</version>
</dependency>
<dependency>
  <groupId>org.apache.flink</groupId>
  <artifactId>flink-streaming-scala_2.11</artifactId>
  <version>1.8.1</version>
</dependency>
<dependency>
  <groupId>org.apache.flink</groupId>
  <artifactId>flink-clients_2.11</artifactId>
  <version>1.8.1</version>
</dependency>
```

Load Model from tf saved model tar, and do prediction

```
class GarbageClassificationInferenceModel(savedModelTarBytes: Array[Byte], inputShape: Array[Int],  
  extends InferenceModel( concurrentNum = 1) with Serializable {  
  doLoadTF(savedModelTarBytes, inputShape, ifReverseInputChannels, meanValues, scale, input)  
  println(this)
```

)

```
class InferenceMapFunction(savedModelTarBytes: Array[  
  extends RichMapFunction[A, B] {  
  var model: GarbageClassificationInferenceModel = _  
  
  override def open(parameters: Configuration): Unit =  
    model = new GarbageClassificationInferenceModel(s  
  }  
  
  override def close(): Unit = {  
    model.release()  
  }  
  
  override def map(in: A): B = {  
    ...  
    val output = model.doPredict(input)  
    ...  
  }  
}
```

)



Apache Flink

THANKS





Apache Flink

关注微信公众号 Vererica

获取 Apache Flink 极客挑战赛最新信息

大赛由 intel 、阿里云计算平台事业部、天池平台联合举办

- 1.两大赛题详细解析
- 2.大赛最新环节通知预告
- 3.优秀项目经验分享
- 4.Flink 系列学习干货

